

École chercheurs MEXICO, La Rochelle, Mars 2018


Model Based Optimization

Nicolas Durrande, nicolas@proowler.io

PROWLER.io, Cambridge (UK) – Mines St-Étienne (France)

Model based optimisation

If the number of function evaluations is limited, we can run the optimization on the model instead of running it directly on the function




In the end, we hope that:

$$\operatorname{argmin}(m) \approx \operatorname{argmin}(f)$$

$$\min(m) \approx \min(f)$$

Overall framework



In practice, it is risky to take decisions based only on the model...

On the other hand, the model can be used to guide us in the search for the optimum.

Contexte

Expériences numériques comme aide à la conception / décision

- Réponse du code de calcul = performance ou coût
- Recherche des paramètres optimaux :

$$x^* = \arg \min \text{cout}(x) \text{ ou } \arg \max \text{perf}(x)$$

- L'optimisation nécessite beaucoup d'appels au code
- Métamodèle : solution naturelle

Lien avec la problématique précédente

Le métamodèle doit être précis seulement dans les régions importantes (proche de l'extremum) \Rightarrow répartition des expériences "ciblée"

Démarche nécessairement séquentielle

Planification d'expériences et optimisation globale

Optimisation locale

Amélioration depuis un point initial

Optimisation globale : le **compromis exploration / intensification**

- Exploration : recherche partout dans l'espace pour ne pas rater la zone optimale
- Intensification : une fois une zone identifiée : on recherche le minimum local

Dans un contexte de planification d'expériences

- Exploration : remplissage d'espace
- Intensification : "ciblage"

Introduction à l'optimisation globale : l'algorithme DIRECT

Garanti sans métamodèle !

DIRECT : Dividing RECTangles


- Découpage de l'espace en (hyper)rectangles
- Un échantillon au centre de chaque rectangle
- On divise les rectangles les plus “intéressants” :
 - ▶ soit les plus grands (exploration)
 - ▶ soit ceux qui ont une valeur au centre basse (intensification)
- Pour diviser : ajout de 2 points, division en 3




D. Jones, C. Perttunen, B. Stuckman (1993)

Lipschitzian optimization without the Lipschitz constant

Journal of Optimization Theory and Applications 79(1), 157-181




Exemple en dimension 2



Après 191 évaluations

- Echantillonnage intense dans la zone de l'optimum
- Bonne exploration



Source figures :



D. E. Finkel

DIRECT Optimization
Algorithm User Guide
(2003)

Intêret et limites

- + Exploration de tout l'espace de recherche
- + Stratégie robuste
- Limité aux petites dimensions
- Exploitation limitée de l'information

⇒ même principe général, avec un métamodèle ?

Optimisation et métamodèle : ce qu'on est tenté de faire...


“Le métamodèle donne l'optimum”

- On cherche le minimum x^* sur le métamodèle
- On évalue le vrai $y(x^*)$ sur le simulateur

⇒ C'est fini !

Répartition de l'effort

- Plan initial : 49 expériences
- 98% exploration, 2% exploitation



Que faire si x^* n'est pas bon ?

Optimisation et métamodèle : ce qu'il faut faire

Si le budget est fixe

- On divise le budget en 2
- Budget 1 : plan initial (LHS)
- Budget 2 : optimisation



Utilisation **séquentielle** du métamodèle

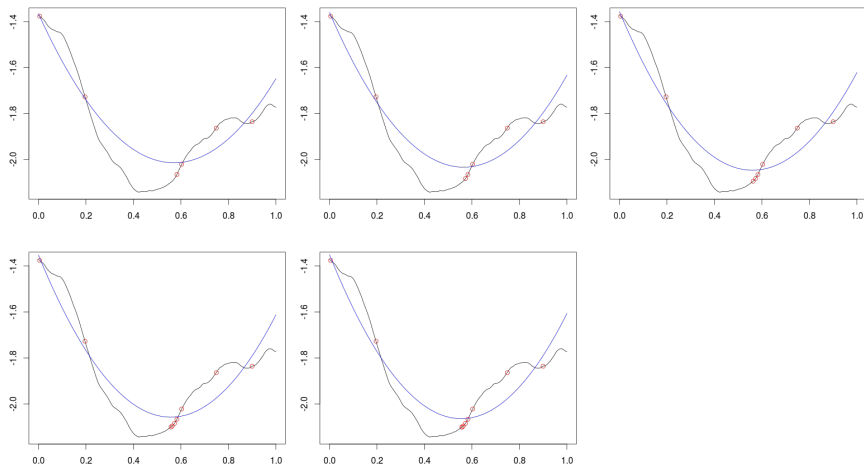
- Métamodèle initial : *a priori* peu précis
- Le métamodèle sert à **choisir** pour les nouvelles observations
- A chaque nouvelle observation : amélioration du métamodèle

Optimisation basée sur les modèles polynomiaux

Principe

- On construit une surface de réponse $y = \beta_0 + \beta_1x + \beta_2x^2$
- On cherche le point qui minimise la surface de réponse
- On ajoute ce point
- On met à jour la surface de réponse
- On recommence...





Optimisation basée sur les modèles polynomiaux

Problème : modèle “rigide”

Le modèle ne s'ajuste pas aux données : $Y = \mathbf{X}\beta + \epsilon$

Pas de convergence vers un modèle précis, même localement

Solutions

1. Augmenter le degré du polynôme
⇒ **risque de surapprentissage & d'instabilité !**
2. Supprimer des points
⇒ méthode **de région de confiance**

Régions de confiance : principe

Modèle quadratique “creux”

- Valide à l'intérieur d'une région de confiance (petite)
- Construit uniquement avec les points à l'intérieur de la région
- Selon les valeurs des simulations, on modifie la taille de la région

Gestion de la région de confiance

A chaque itération :

- $\hat{y}(x^*)$ bon \Rightarrow confiance dans le modèle : on augmente la taille
- $\hat{y}(x^*)$ mauvais \Rightarrow modèle peu fiable : on diminue la taille

+ beaucoup de règles pour sélectionner les points et enrichir le plan d'expériences

Illustration (source : F. Vanden Berghen)




Illustration (source : F. Vanden Berghen)





Illustration (source : F. Vanden Berghen)



Avantages et inconvénients

Avantages

- Garantie de convergence
- Méthodes assez parcimonieuses
- Robuste
- Accepte un assez grand nombre de variables



Conn, Scheinberg, and Vicente

Introduction to derivative-free optimization
MPS-SIAM Series on Optimization (2009)


Méthode locale

- Pas de métamodèle final utilisable globalement
- Proche des méthodes de gradient


Global optimization methods are a trade-off between

- Exploitation of past good results
- Exploration of the space

How can GPR models be helpful?



In our example, the best observed value is 1.79




Various criteria can be studied

- probability of improvement
- Expected improvement

Probability of Improvement:

$$PI(x) = cdf \left(\frac{\min(F) - m(x)}{\sqrt{c(x, x)}} \right)$$



The point with the highest PI is often very close to the best observed value. We can show that there is a x in the neighbourhood of x^* such that $PI(x) \geq 0.5$.


For such points, the improvement cannot be large...

Can we find another criterion?

Expected Improvement:


$$EI(x) = \int_{-\infty}^{\min(F)} \max(0, Y(x)) \, dy(x) = \dots =$$

$$\sqrt{c(x, x)}(u(x)cdf(u(x)) + pdf(u(x))) \quad \text{with } u(x) = \frac{\min(F) - m(x)}{\sqrt{(c(x, x))}}$$




Expected Improvement

Let's see how it works... iteration 1




Expected Improvement

Let's see how it works... iteration 2




Expected Improvement

Let's see how it works... iteration 3




Expected Improvement

Let's see how it works... iteration 4



Expected Improvement

Let's see how it works... iteration 5



Expected Improvement

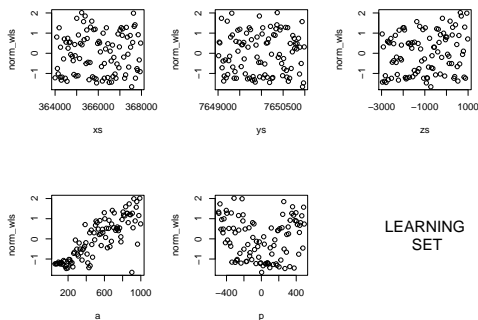
This algorithm is called **Efficient Global Optimization** (EGO, Jones et al., 1998):

1. make an initial design of experiments X and calculate the associated F , $t = \text{length}(F)$
 2. built a GP from (X, F) (max. log-likelihood on σ and θ_i 's)
 3. $X_{t+1} = \arg \max_x EI(x)$
 4. calculate $F_{t+1} = f(X_{t+1})$, increment t
 5. stop ($t > t^{\max}$) or go to 2.
-
- + EGO provides a good trade-off between exploitation and exploration without arbitrary parameters.
 - + It requires few function observations (10 in the example) to get close to optimal regions.

Example in 5d: surface displacements misfit minimization

⇒ demo with `mainInversionPunctualDisplSource.R`

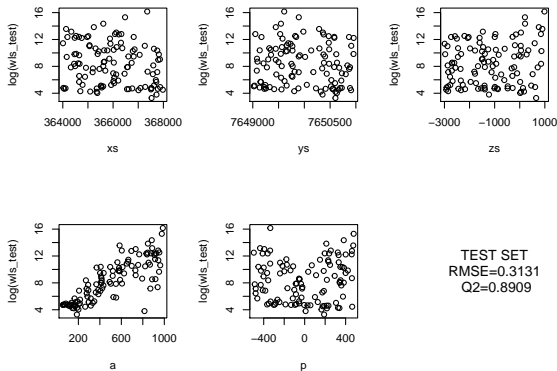
!!! normalize the data: WLS has a few very large values, it is always > 0 : make it more gaussian, $wls_norm = \log(1 + wls)$ and all x 's and wls_norm between 0 and 1.



LEARNING
SET


100 $\{x_s, y_s, z_s, a, p\}$
points chosen through
an optimized Latin
Hypercube Sampling
(R libraries
DiceDesign or lhs).

(demo with `mainInversionPunctualDisplSource.R`, cont.)



110 random $\{xs, ys, zs, a, p\}$ test points.


(demo with mainInversionPunctualDisplSource.R, cont.)



(demo with `mainInversionPunctualDisplSource.R`, cont.)

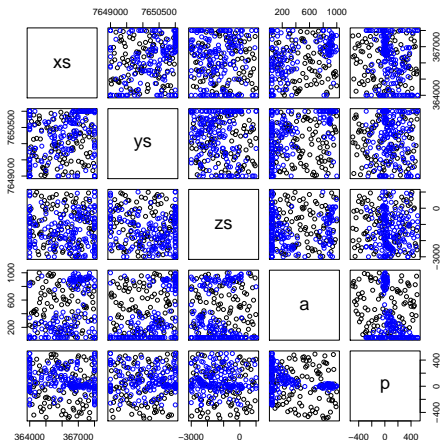
EGO parameters: anisotropic Matérn 5/2 kernel, GP updated

(log-likelihood maximized) every 5 added points, BFGS with bounded variables (from `optim()` function) restarted from random initial points for maximizing log-likelihood and EI.



Preferential sampling of good regions of S , but global therefore sometimes increasing WLS. Lower bound on θ_i 's increased from 0.08 to 0.1 at $t = 250$ (x_i 's and θ_i 's normed between 0 and 1).

(demo with mainInversionPunctualDisplSource.R, cont.)

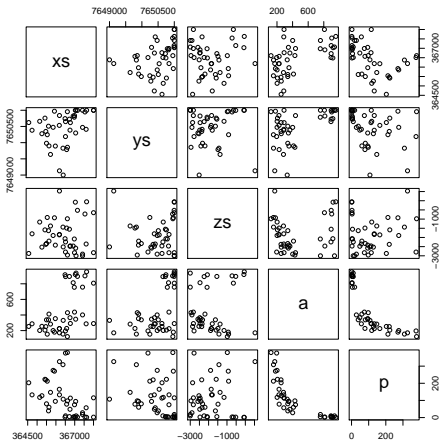


Black: LHS initial points.

Blue: EGO points.


Note the patterns in new points. Accumulation at lower bound of a and mid interval of p before $t = 250$.

(demo with mainInversionPunctualDisplSource.R, cont.)



15% best sampled points. Note the “function” for the (a, p) pair, i.e., $a^*(p^*)$.

(demo with mainInversionPunctualDisplSource.R, cont.)



Mogi model only dependency
in a and p is through $a^3 \times p$:
it is not identifiable.

EGO tells it by preferential
sampling in the valley

$$a^3 \times p = \text{const.} = a^{*3} \times p^*$$

Other EGO output: a statistical model of WLS.

The last length scales are an indication of the sensitivity of WLS to each variable: a , p and zs are very sensitive (θ_i 's small, in $[0.08, 0.1]$), xs a little sensitive (θ in $[0.1, 2.5]$) and ys insensitive ($\theta \approx 3$).

Difficulties and challenges with EGO

- Standard GPs are limited to $n \approx 1000$ points (covariance matrix inversion).
- EGO clusters points in good regions, the covariance matrix may become ill-conditioned if length scales θ_i are too large w.r.t. X .
- Although the method perfectly applies to large dimensional spaces ($d > 100$), larger d may require larger n , back 2 lines above.
- EGO does not converge in the traditional sense: it creates dense samples in the volume of S . The efficiency comes from the order in which points are sampled.

⇒ these are the topics of current research. Let's mention a few extensions next.

EGO continuations

- Parallelized EGO: estimate the El of groups of points, cf. Ginsbourger et al.
- Finite budget: El of a single x is only optimal at the last iteration. Theory of dynamic El , cf. Ginsbourger et al.
- EGO and bad covariance matrix conditioning: replace points that are close-by by one point and the associated derivatives (cf. M. Osborn, L. Laurent), regularizations (cf. Le Riche et al.)
- SUR strategies: (Step-wise Uncertainty Reduction), reduce the entropy of the optimum (cf. Vasquez et al.), or the average probability of incursions below $\min(F)$ (cf. Picheny).

Related problems addressed with GPs

- EGO with constraints: $\min_x f(x)$ s.t. $g(x) \leq 0$, multiply the EI by the probability of constraints satisfaction.
- GP for target attainment: find the set of x s.t. $f(x) = T$, change the EI into $c(x, x) \times \text{pdf}((T - m(x))/\text{sqrt}(c(x, x)))$, cf. Picheny et al.
- GP for probability estimation: find $\mathbb{P}(f(x, U) \leq T)$ where U is a random vector.
- GP for multi-objective optimization: $\min_x \{f_1(x), \dots, f_m(x)\}$, cf. Binois et al.

Robust optimization

Can EGO be adapted when observations are noisy?


First of all, using the current best observation as a minimum does not make much sense...

Some solutions are

- S1 Build a new model that interpolates $m(X)$ at X where $m(X)$ accounts for the noise (non interpolating GP, e.g. with a white noise part in the kernel).
- S2 Include observation noise and replace $\min(F)$ by $\min(m(X))$ in the EI expression
- S3 Similar to 2 but consider an Expected Mean Improvement (V. Picheny).

Solution 1


iteration 0



(noisy observations and their denoised versions are both shown as black crosses)

Solution 1


iteration 1



(noisy observations and their denoised versions are both shown as black crosses)

Solution 1


iteration 2



(noisy observations and their denoised versions are both shown as black crosses)

Solution 1


iteration 3



(noisy observations and their denoised versions are both shown as black crosses)

Solution 1

iteration 4



(noisy observations and their denoised versions are both shown as black crosses)

Kernel Design

Making new from old: Many operations can be applied to psd functions while retaining this property

Kernels can be:

- Summed together

- ▶ On the same space $k(x, y) = k_1(x, y) + k_2(x, y)$
- ▶ On the tensor space $k(x, y) = k_1(x_1, y_1) + k_2(x_2, y_2)$

- Multiplied together

- ▶ On the same space $k(x, y) = k_1(x, y) \times k_2(x, y)$
- ▶ On the tensor space $k(x, y) = k_1(x_1, y_1) \times k_2(x_2, y_2)$

- Composed with a function


- ▶ $k(x, y) = k_1(f(x), f(y))$

How can this be useful?

Sum of kernels over the same space

Example (The Mauna Loa observatory dataset)

This famous dataset compiles the monthly CO_2 concentration in Hawaii since 1958.





Let's try to predict the concentration for the next 20 years.

Sum of kernels over the same space

We first consider a squared-exponential kernel:

$$k_{se}(x, y) = \sigma^2 \exp\left(-\frac{(x - y)^2}{\theta^2}\right)$$



The results are terrible!

Sum of kernels over the same space


What happen if we sum both kernels?

$$k(x, y) = k_{se1}(x, y) + k_{se2}(x, y)$$

Sum of kernels over the same space

What happen if we sum both kernels?

$$k(x, y) = k_{se1}(x, y) + k_{se2}(x, y)$$



The model is drastically improved!

Sum of kernels over the same space


We can try the following kernel:

$$k(x, y) = \sigma_0^2 x^2 y^2 + k_{se1}(x, y) + k_{se2}(x, y) + k_{per}(x, y)$$

Sum of kernels over the same space

We can try the following kernel:

$$k(x, y) = \sigma_0^2 x^2 y^2 + k_{se1}(x, y) + k_{se2}(x, y) + k_{per}(x, y)$$




Once again, the model is significantly improved.

Sum of kernels over tensor space

Property

$$k(x, y) = k_1(x_1, y_1) + k_2(x_2, y_2)$$

is valid covariance structure.






Remark: From a GP point of view, k is the kernel of

$$Z(x) = Z_1(x_1) + Z_2(x_2)$$

Sum of kernels over tensor space

We can have a look at a few sample paths from Z :



⇒ They are additive (up to a modification)

Tensor Additive kernels are very useful for

- Approximating additive functions
- Building models over high dimensional inputs spaces

Product over the same space


Property

$$k(x, y) = k_1(x, y) \times k_2(x, y)$$

is valid covariance structure.

Example

We consider the product of a squared exponential with a cosine:



Product over the tensor space


Property

$$k(x, y) = k_1(x_1, y_1) \times k_2(x_2, y_2)$$

is valid covariance structure.

Example

We multiply 2 squared exponential kernel



Calculation shows this is the usual 2D squared exponential kernel.

Composition with a function

Property

Let k_1 be a kernel over $D_1 \times D_1$ and f be an arbitrary function $D \rightarrow D_1$, then

$$k(x, y) = k_1(f(x), f(y))$$

is a kernel over $D \times D$.

proof

$$\sum \sum a_i a_j k(x_i, x_j) = \sum \sum a_i a_j k_1(\underbrace{f(x_i)}_{y_i}, \underbrace{f(x_j)}_{y_j}) \geq 0$$

Remarks:



- k corresponds to the covariance of $Z(x) = Z_1(f(x))$
- This can be seen as a (non-linear) rescaling of the input space

Example

We consider $f(x) = \frac{1}{x}$ and a Matérn 3/2 kernel

$$k_1(x, y) = (1 + |x - y|)e^{-|x - y|}.$$

We obtain:





All these transformations can be combined!


Example

$k(x, y) = f(x)f(y)k_1(x, y)$ is a valid kernel.

This can be illustrated with $f(x) = \frac{1}{x}$ and
 $k_1(x, y) = (1 + |x - y|)e^{-|x-y|}$:



Example



⇒ R demo

Other kernel design methods

There are two other popular methods for kernel design:

- Bochner Theorem

There is an equivalence between positive measures and stationnary positive definite functions.

- Linear operators

If the function to approximate has particular properties that can be obtained via a linear transform, it is possible to build a GP with the wanted properties. For example, one can build symmetric GPs or GPs with integral equal to zero.